

COMS 4701 Artificial Intelligence

Homework 2 – Due date: February 23rd, 2020

Justify all your work to receive full credit

I- Written


1. Local Search

Provide three advantages and three disadvantages of using local search algorithms. Explain.

2. Local Search

The 6-queen problem requires that you place 6 queens on a 6 × 6 chessboard, so as no queen attacks another one located in the same row, column or diagonal. We define the states to be any configuration of the board where we have one queen per column. The successor function moves a single queen to another square in the same column. For a state S, consider the evaluation function $\text{fitness}(S)$, as the number of non-attacking pairs of queens in state S.

- (a) How many possible states are there in total?
- (b) How many successor states are there for each state?
- (c) What is the fitness function for the following state:



- (d) We would like to use a genetic algorithm, where each state is represented by a sequence of 6 digits. Each digit represents the position of the queen in that column. For instance, the board above would be represented by the sequence **3 5 1 5 4 3**.

What genetic operator should we use to implement the successor function described in this problem? Consider only the operators of cross-over, selection and mutation.

3. Association Rules

Consider the following items and transactions.

Item id	Item
1	diapers
2	beer
3	water
4	coffee
5	milk

TID	Transaction
100	1 2 4 5
200	2 5
300	1 2 4 5
400	1 2 5
500	1 2 3 4 5
600	3 4 5

- (a) For a minimum support of 50%, use the Apriori algorithm to find all frequent itemsets in the transaction table.
- (b) For a minimum confidence of 80%, use the Apriori algorithm to find all strong association rules (report support and confidence) of the form:

$$Item1 \rightarrow Item2 \text{ (support, confidence)}$$

4. Constraint Satisfaction Problems

We are going to consider a simplified version of the sudoku puzzle:

4	<i>A</i>	<i>B</i>	<i>C</i>
<i>D</i>	<i>E</i>	<i>F</i>	4
<i>G</i>	2	<i>H</i>	<i>I</i>
<i>J</i>	<i>K</i>	3	<i>L</i>

The solution to this problem is to generally assign values between 1 and 4 to each of the letters while satisfying the condition; each number (1-4) must appear exactly once in each row, column, and 2x2 box, without changing the numerals that are already in the grid.

- (a) What is the size of the state space for this CSP? You do not have to consider constraints.
- (b) List all unary constraints
- (c) If we choose to assign the first variable using the minimum remaining values (MRV) heuristic, which variable(s) would get chosen? List all such variables if there is a tie.
- (d) Prior to assigning any variables, use both unary constraints and arc consistency due to binary constraints to reduce the domains of the problem.

II- Programming

Read Carefully:

- You must name your file `puzzle.py`. Any submission that does not follow this naming will not be graded and will receive a zero.
- A skeleton of each function has been provided to you in `puzzle.py`. You are expected to **ONLY** write code in the functions and blocks specified. In any case, **DO NOT** modify the function signatures, or any code that is not specified to be modifiable for any reason. This will be run by our autograder, so any unexpected modifications that make it malfunction will receive a zero. **You can, however, modify the main function** as necessary to make the program run as specified in the assignment. Test cases have been provided in the main function in the skeleton code.
- Only Python 3.x versions will be graded.
- To receive points, make sure your code runs. We recommend using Spyder, Pycharm or Google Colab. They all allow you to download .py files. Be aware that if you write your code in some platforms like Codio and copy and paste it in a text file, there may be spurious characters in your code, and your code will not compile. Always ensure that your .py compiles. Code that does not run will not receive any points.
- **Submission Instructions:** They will be communicated to you over the course of the week.

In this assignment you will create an agent to solve the N-puzzle game. You will implement and compare several search algorithms, and collect some statistics related to their performances. Visit mypuzzle.org/sliding for the game's rules. Please read all sections carefully:

I. Introduction

II. Algorithm Review

III. What You Need To Submit
IV. What Your Program Outputs
V. Implementation and Testing
VI. Before You Finish

I. Introduction

The N-puzzle game consists of a board holding $N = m^2 - 1$ distinct movable tiles, plus one empty space. There is one tile for each number in the set $\{0, 1, \dots, m^2 - 1\}$. In this assignment, we will represent the blank space with the number 0 and focus on the $m = 3$ case (8-puzzle).

In this combinatorial search problem, the aim is to get from any initial board state to the configuration with all tiles arranged in ascending order $\{0, 1, \dots, m^2 - 1\}$ – this is your goal state. The search space is the set of all possible states reachable from the initial state. Each move consists of swapping the empty space with a component in one of the four directions {'Up', 'Down', 'Left', 'Right'}. Give each move a cost of one. Thus, the total cost of a path will be equal to the number of moves made.

II. Algorithm Review

Recall from lecture that search begins by visiting the root node of the search tree, given by the initial state. Three main events occur when visiting a node:

- First, we remove a node from the frontier set.
- Second, we check if this node matches the goal state.
- If not, we then expand the node. To expand a node, we generate all of its immediate successors and add them to the frontier, if they (i) are not yet already in the frontier, and (ii) have not been visited yet.

This describes the life cycle of a visit, and is the basic order of operations for search agents in this assignment—(1) remove, (2) check, and (3) expand. We will implement the assignment algorithms as described here. Please refer to lecture notes for further details, and review the lecture pseudo-code before you begin.

IMPORTANT: You may encounter implementations that attempt to short-circuit this order by performing the goal-check on successor nodes immediately upon expansion of a parent node. For example, Russell Norvig's implementation of BFS does precisely this. Doing so may lead to edge-case gains in efficiency, but do not alter the general characteristics of complexity and optimality for each method. For simplicity and grading purposes in this assignment, do not make such modifications to algorithms learned in lecture.

III. What You Need To Submit

Your job in this assignment is to write `puzzle.py`, which solves any 8-puzzle board when given an arbitrary starting configuration. The program will be executed as follows:

```
$ python3 puzzle.py <method><board>
```

The method argument will be one of the following. You must implement all three of them:

`bfs` (Breadth-First Search)
`dfs` (Depth-First Search)
`ast` (A-Star Search)

The board argument will be a comma-separated list of integers containing no spaces. For example, to use the bread-first search strategy to solve the input board given by the starting configuration $\{0, 8, 7, 6, 5, 4, 3, 2, 1\}$, the

program will be executed like so (with no spaces between commas):

```
$ python3 puzzle.py bfs 0,8,7,6,5,4,3,2,1
```

IV. What Your Program Outputs

Your program will create and/or write to a file called `output.txt`, containing the following statistics:

`path_to_goal`: the sequence of moves taken to reach the goal

`cost_of_path`: the number of moves taken to reach the goal

`nodes_expanded`: the number of nodes that have been expanded

`search_depth`: the depth within the search tree when the goal node is found

`max_search_depth`: the maximum depth of the search tree in the lifetime of the algorithm

`running_time`: the total running time of the search instance, reported in seconds

`max_ram_usage`: the maximum RAM usage in the lifetime of the process as measured by the `ru_maxrss` attribute in the `resource` module, reported in megabytes

Example 1: Breadth-First Search

Suppose the program is executed for breadth-first search as follows:

```
$ python3 puzzle.py bfs 1,2,5,3,4,0,6,7,8
```

This should result in the solution path:

$parent = \begin{array}{|c|c|c|} \hline 1 & 2 & 5 \\ \hline 3 & 4 & \\ \hline 6 & 7 & 8 \\ \hline \end{array} \Rightarrow child = \begin{array}{|c|c|c|} \hline 1 & 2 & \\ \hline 3 & 4 & 5 \\ \hline 6 & 7 & 8 \\ \hline \end{array}$

$parent = \begin{array}{|c|c|c|} \hline 1 & 2 & \\ \hline 3 & 4 & 5 \\ \hline 6 & 7 & 8 \\ \hline \end{array} \Rightarrow child = \begin{array}{|c|c|c|} \hline 1 & & 2 \\ \hline 3 & 4 & 5 \\ \hline 6 & 7 & 8 \\ \hline \end{array}$

$parent = \begin{array}{|c|c|c|} \hline 1 & & 2 \\ \hline 3 & 4 & 5 \\ \hline 6 & 7 & 8 \\ \hline \end{array} \Rightarrow child = \begin{array}{|c|c|c|} \hline & 1 & 2 \\ \hline 3 & 4 & 5 \\ \hline 6 & 7 & 8 \\ \hline \end{array}$

The output file will contain exactly the following lines:

`path_to_goal`: ['Up', 'Left', 'Left']

`cost_of_path`: 3

`nodes_expanded`: 10

`search_depth`: 3

`max_search_depth`: 4

`running_time`: 0.00188088

`max_ram_usage`: 0.07812500

Example 2: Depth-First Search

Suppose the program is executed for depth-first search as follows:

```
$ python3 puzzle.py dfs 1,2,5,3,4,0,6,7,8
```

This should result in the solution path:

$parent = \begin{array}{|c|c|c|} \hline 1 & 2 & 5 \\ \hline 3 & 4 & \\ \hline 6 & 7 & 8 \\ \hline \end{array} \Rightarrow child = \begin{array}{|c|c|c|} \hline 1 & 2 & \\ \hline 3 & 4 & 5 \\ \hline 6 & 7 & 8 \\ \hline \end{array}$

$parent = \begin{array}{|c|c|c|} \hline 1 & 2 & \\ \hline 3 & 4 & 5 \\ \hline 6 & 7 & 8 \\ \hline \end{array} \Rightarrow child = \begin{array}{|c|c|c|} \hline 1 & & 2 \\ \hline 3 & 4 & 5 \\ \hline 6 & 7 & 8 \\ \hline \end{array}$

$parent = \begin{array}{|c|c|c|} \hline 1 & & 2 \\ \hline 3 & 4 & 5 \\ \hline 6 & 7 & 8 \\ \hline \end{array} \Rightarrow child = \begin{array}{|c|c|c|} \hline & 1 & 2 \\ \hline 3 & 4 & 5 \\ \hline 6 & 7 & 8 \\ \hline \end{array}$

The output file will contain exactly the following lines:

```
path_to_goal: ['Up', 'Left', 'Left']
```

```
cost_of_path: 3
```

```
nodes_expanded: 181437
```

```
search_depth: 3
```

```
max_search_depth: 66125
```

```
running_time: 5.01608433
```

```
max_ram_usage: 4.23940217
```

More test cases are provided in **the FAQs**.

Note on Correctness

All variables, except `running_time` and `max_ram_usage`, have **one and only one** correct answer when running BFS and DFS. A* `nodes_expanded` might vary depending on implementation details. You'll be fine as long as your algorithm follows all specifications listed in these instructions.

As `running_time` and `max_ram_usage` values vary greatly depending on your machine and implementation details, there is no "correct" value to look for. They are for you to monitor time and space complexity of your code, which we highly recommend. A good way to check the correctness of your program is to walk through small examples by hand, like the ones above.

V. Implementation and Testing

For your first programming project, we are providing hints and explicit instructions. Before posting a question on the discussion board, make sure your question is not already answered here or in the FAQs.

1. Implementation

You will implement the following three algorithms as demonstrated in lecture. In particular:

- **Breadth-First Search.** Use an explicit queue, as shown in lecture.
- **Depth-First Search.** Use an explicit stack, as shown in lecture.
- **A-Star Search.** Use a priority queue, as shown in lecture. For the choice of heuristic, use the Manhattan priority function that is, the sum of the distances of the tiles from their goal positions. Note that the blanks space is not considered an actual tile here.

2. Order of Visits

In this assignment, where an arbitrary choice must be made, we always **visit** child nodes in the “UDLR” order; that is, [‘Up’, ‘Down’, ‘Left’, ‘Right’] in that exact order. Specifically:

- **Breadth-First Search.** Enqueue in UDLR order; de-queuing results in UDLR order.
- **Depth-First Search.** Push onto the stack in reverse-UDLR order; popping off results in UDLR order.
- **A-Star Search.** Since you are using a priority queue, what happens with duplicate keys? How do you ensure nodes are retrieved from the priority queue in the desired order?

3. Submission Test Cases

Run all three of your algorithms on the following test cases:

Test Case 1

```
$python3 puzzle.py bfs 3,1,2,0,4,5,6,7,8
```

```
$python3 puzzle.py dfs 3,1,2,0,4,5,6,7,8
```

```
$python3 puzzle.py ast 3,1,2,0,4,5,6,7,8
```

Test Case 2

```
$python3 puzzle.py bfs 1,2,5,3,4,0,6,7,8
```

```
$python3 puzzle.py dfs 1,2,5,3,4,0,6,7,8
```

```
$python3 puzzle.py ast 1,2,5,3,4,0,6,7,8
```

Make sure your code passes at least these test cases and follows our formatting exactly. The results of each test are assessed by 8 items: 7 are listed in **Section IV. What Your Program Outputs**. The last point is for code that executes and produces any output at all. Each item is worth 0.75 point.

4. Grading and Stress Tests

We will grade your project by running **additional test cases** on your code. In particular, there will be five test cases in total, each tested on all three of your algorithms, for a total of **15** distinct tests. Similar to the submission test cases, each test will be graded by 8 items, for a total of 90 points. **Plus, we give 10 points for code completing all 15 test cases within 10 minutes.** If you implement your code with reasonable designs of data structures, your code will solve all 15 test cases within a minute in total. We will be using a wide variety of inputs to stress-test your algorithms to check for correctness of implementation. So, we recommend that you test your own code extensively.

Don’t worry about checking for malformed input boards, including boards of non-square dimensions, other sizes, or unsolvable boards.

You will not be graded on the absolute values of your running time or RAM usage statistics. The values of these statistics can vary widely depending on the machine. **However, we recommend that you take advantage of them in testing your code.** Try batch-running your algorithms on various inputs, and plotting your results on a graph to learn more about the space and time complexity characteristics of your code. Just because an algorithm provides the correct path to goal does not mean it has been implemented correctly.

5. Tips on Getting Started

Begin by writing a class to represent the **state** of the game at a given turn, including parent and child nodes. We suggest writing a separate **solver** class to work with the state class. Feel free to experiment with your design, for example including a **board** class to represent the low-level physical configuration of the tiles, delegating the high-level functionality to the state class.

You will not be graded on your design, so you are at a liberty to choose among your favorite programming paradigms. Students have successfully completed this project using an entirely object-oriented approach, and

others have done so with a purely functional approach. Your submission will receive full credit as long as your puzzle program outputs the correct information.

VI. Before You Finish

- **Make sure** your code passes at least the submission test cases.
- **Make sure** your algorithms generate the correct solution for an arbitrary solvable problem instance of 8-puzzle.
- **Make sure** your program always terminates without error, and in a reasonable amount of time. **You will receive zero points from the grader if your program fails to terminate. Running times of more than a minute or two may indicate a problem with your implementation.** If your implementation exceeds the time limit allocated (20 minutes for all test cases), your grade may be incomplete.
- **Make sure** your program output follows the specified format exactly. In particular, for the path to goal, use square brackets to surround the list of items, use single quotes around each item, and capitalize the first letter of each item. Round floating-point numbers to 8 places after the decimal. You will not receive proper credit from the grader if your format differs from the provided examples above.